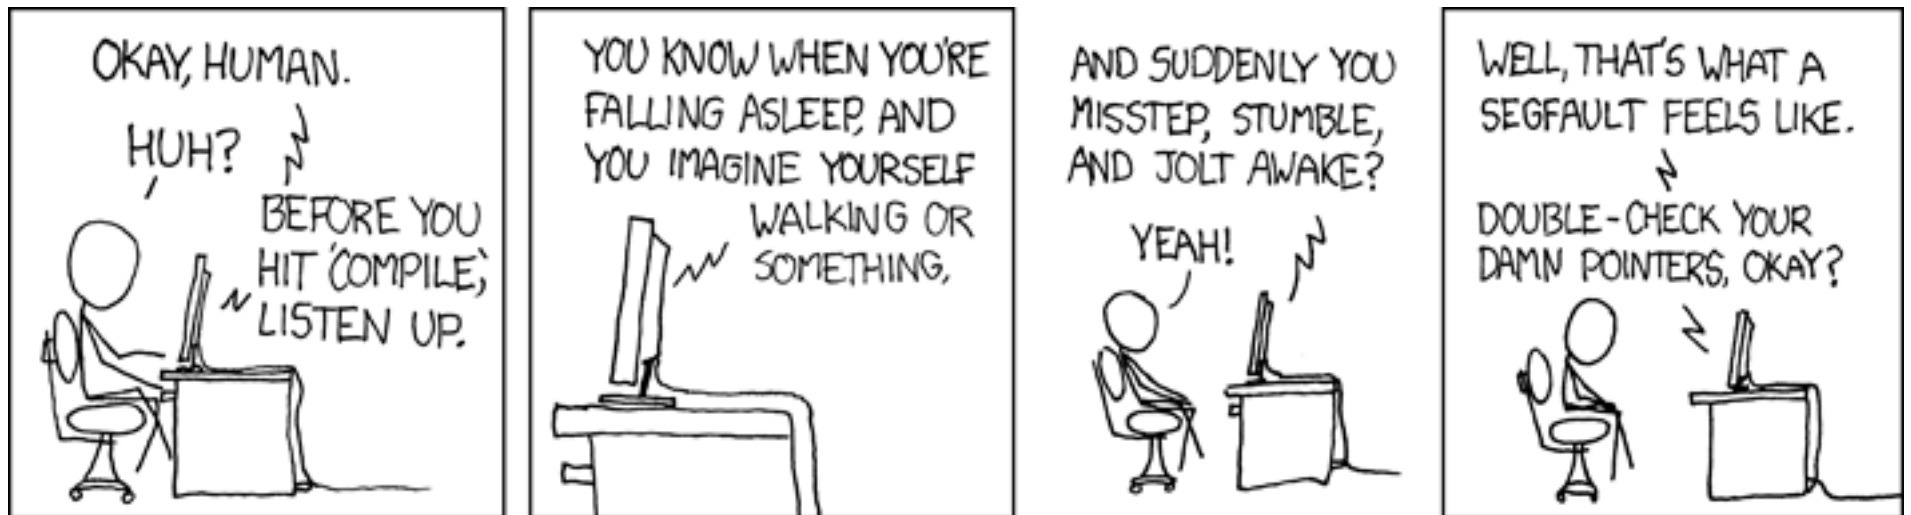


# Goals for today


- The little button that wouldn't :(
  - the `volatile` keyword
- How do functions work in C?
- Management of runtime stack, register use
- Multiple loads and stores: `ldm` and `stm`



**button.c**

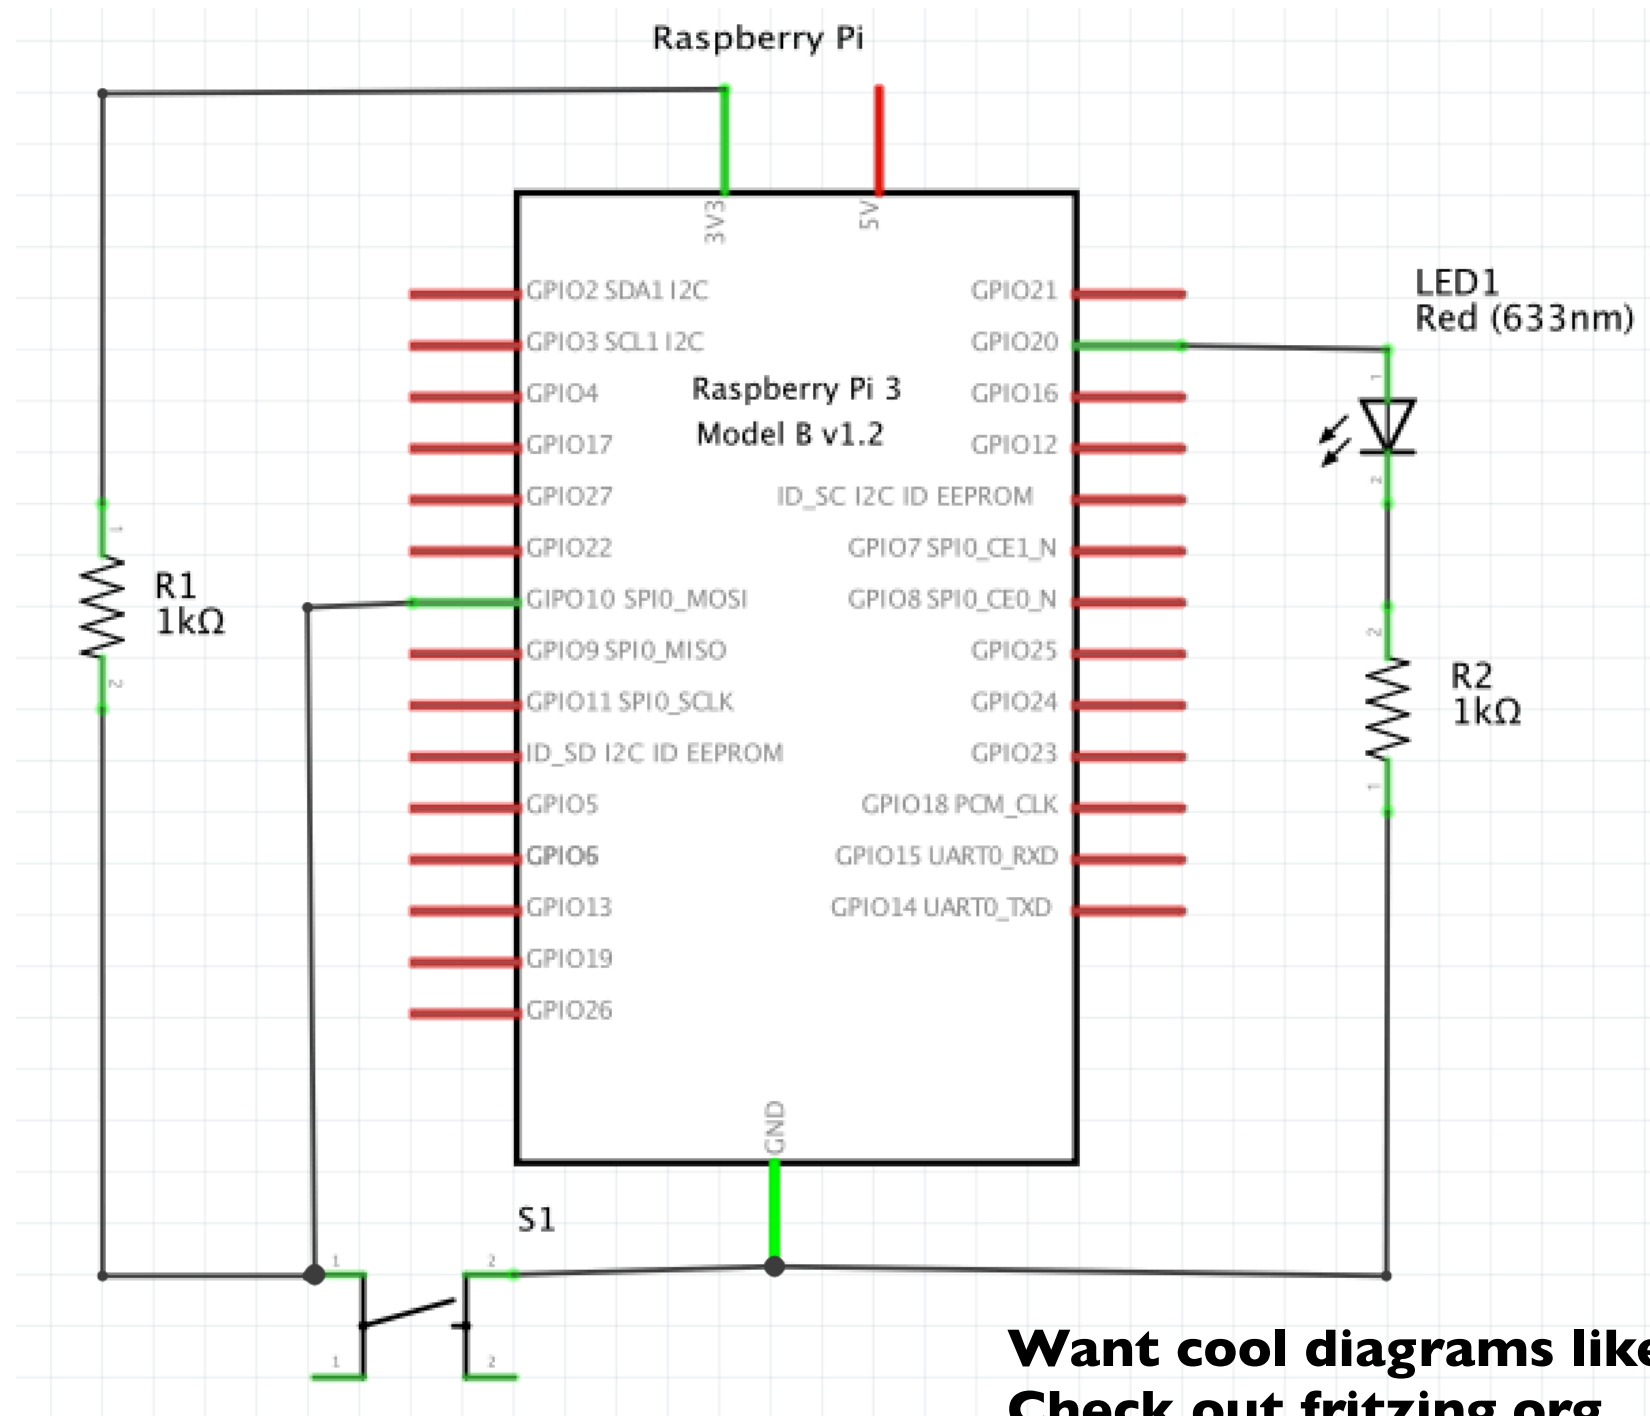
**The little button that wouldn't**

# button.c: The little button that wouldn't



**Want cool diagrams like this?  
Check out [fritzing.org](http://fritzing.org)**

# button.c: The little button that wouldn't



**Want cool diagrams like this?  
Check out [fritzing.org](http://fritzing.org)**

# button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20

unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0  = (unsigned int *)0x2020001C;
unsigned int * const CLRO  = (unsigned int *)0x20200028;
unsigned int * const LEV0  = (unsigned int *)0x20200034;

void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SET0 = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEV0 & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLRO = 1 << 20;
    }
}
```

# button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20
```

```
unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0  = (unsigned int *)0x2020001C;
unsigned int * const CLRO  = (unsigned int *)0x20200028;
unsigned int * const LEV0  = (unsigned int *)0x20200034;
```

```
void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SET0 = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEV0 & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLRO = 1 << 20;

    }
}
```

## Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
    0:    ldr    r3, [pc, #28] ; 24 <main+0x24>
    4:    ldr    r0, [r3, #52] ; 0x34
    8:    mov    r1, #0
    c:    mov    r2, #1
   10:    tst    r0, #1024 ; 0x400
   14:    stmib   r3, {r1, r2}
   18:    bne     20 <main+0x20>
   1c:    b       1c <main+0x1c>
   20:    b       20 <main+0x20>
   24:    .word    0x20200000
```

# button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
// released and turns off GPIO 20
```

```
unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0  = (unsigned int *)0x2020001C;
unsigned int * const CLRO  = (unsigned int *)0x20200028;
unsigned int * const LEV0  = (unsigned int *)0x20200034;
```

```
void main(void)
{
```

```
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output
```

```
    while (1) {
```

```
        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;
```

```
        // set GPIO 20 high
```

```
        *SET0 = 1 << 20;
```

```
        // wait until GPIO 10 is high (button release)
```

```
        while ((*LEV0 & (1 << 10)) == 0) ;
```

```
        // clear GPIO 20
```

```
        *CLRO = 1 << 20;
```

```
    }
```

```
}
```

## Compiling with -O2:

Disassembly of section .text.startup:

00000000 <main>:

```
0:   ldr    r3, [pc, #28] ; 24 <main+0x24>
4:   ldr    r0, [r3, #52] ; 0x34
8:   mov    r1, #0
c:   mov    r2, #1
10:  tst    r0, #1024 ; 0x400
14:  stmib  r3, {r1, r2}
18:  bne    20 <main+0x20>
1c:  b      1c <main+0x1c> ?
20:  b      20 <main+0x20> ?
24:  .word   0x20200000 ?
```

# button.c: The little button that wouldn't

```
// This program waits until a button is pressed (GPIO 10)
// and turns on GPIO 20, then waits until the button is
//released and turns off GPIO 20
```

```
unsigned int * const FSEL1 = (unsigned int *)0x20200004;
unsigned int * const FSEL2 = (unsigned int *)0x20200008;
unsigned int * const SET0  = (unsigned int *)0x2020001C;
unsigned int * const CLRO  = (unsigned int *)0x20200028;
unsigned int * const LEV0  = (unsigned int *)0x20200034;
```

```
void main(void)
{
    *FSEL1 = 0; // configure GPIO 10 as input
    *FSEL2 = 1; // configure GPIO 20 as output

    while (1) {

        // wait until GPIO 10 is low (button press)
        while ((*LEV0 & (1 << 10)) != 0) ;

        // set GPIO 20 high
        *SET0 = 1 << 20;

        // wait until GPIO 10 is high (button release)
        while ((*LEV0 & (1 << 10)) == 0) ;

        // clear GPIO 20
        *CLRO = 1 << 20;

    }
}
```

## Compiling with -O2:

Disassembly of section .text.startup:

```
00000000 <main>:
    0:    ldr    r3, [pc, #28] ; 24 <main+0x24>
    4:    ldr    r0, [r3, #52] ; 0x34
    8:    mov    r1, #0
    c:    mov    r2, #1
   10:    tst    r0, #1024 ; 0x400
   14:    stmib   r3, {r1, r2}
   18:    bne     20 <main+0x20>
   1c:    b       1c <main+0x1c>
   20:    b       20 <main+0x20>
   24:    .word    0x20200000
```

## What happened to our testing loops??



# Peripheral Registers

These registers are mapped into the address space of the processor (memory-mapped IO).

These registers may behave differently than memory.

For example: Writing a 1 into a bit in a SET register causes 1 to be output; writing a 0 into a bit in SET register does not affect the output value. Writing a 1 to the CLR register, sets the output to 0; write a 0 to a clear register has no effect. Neither SET or CLR can be read. To read the current value use the LEV (level) register.

# volatile

For an ordinary variable, the compiler can use its knowledge of when it is read/written to optimize accesses as long as it keeps the same externally visible behavior.

However, for a variable that can be read/written externally (by another process, by peripheral), these optimizations will not be valid.

The **volatile** qualifier applied to a variable informs the compiler that it cannot remove, coalesce, cache, or reorder references. The generated assembly must faithfully execute each access to the variable as given in the C code.

## button.c: The little button that **could**

Because we have GPIO pins on the Raspberry Pi, we need to give hints to the C compiler to not optimize out pin reads — they can change externally to the program!

So, we use the `volatile` keyword in front of hardware addresses to do this:

```
volatile unsigned int * const FSEL1 = (unsigned int *)0x20200004;  
volatile unsigned int * const FSEL2 = (unsigned int *)0x20200008;  
volatile unsigned int * const SET0  = (unsigned int *)0x2020001C;  
volatile unsigned int * const CLR0  = (unsigned int *)0x20200028;  
volatile unsigned int * const LEV0  = (unsigned int *)0x20200034;
```

## button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:      file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
    0: e3a00000  mov r0, #0
    4: e12ffff1  bx  lr
```

## button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (int i=0; i < DELAY; i++);

    return 0;
}
```

```
$ objdump -d testLoop.o

testLoop.o:      file format elf32-littlearm

Disassembly of section .text.startup:

00000000 <main>:
    0: e3a00000    mov r0, #0
    4: e12ffffe    bx  lr
```

**No loop — it has been optimized out!**

# button.c: The little button that **could**

There are other times to use volatile, too — delays have a similar problem:

```
#define DELAY 500000000

int main()
{
    for (volatile int i=0; i < DELAY; i++);

    return 0;
}
```

Disassembly of section .text.startup:

```
00000000 <main>:
 0: e24dd008    sub    sp, sp, #8
 4: e3a03000    mov    r3, #0
 8: e58d3004    str    r3, [sp, #4]
 c: e59d3004    ldr    r3, [sp, #4]
10: e59f2028    ldr    r2, [pc, #40]    ; 40 <main+0x40>
14: e1530002    cmp    r3, r2
18: ca000005    bgt    34 <main+0x34>
1c: e59d3004    ldr    r3, [sp, #4]
20: e2833001    add    r3, r3, #1
24: e58d3004    str    r3, [sp, #4]
28: e59d3004    ldr    r3, [sp, #4]
2c: e1530002    cmp    r3, r2
30: daffffff9    ble    1c <main+0x1c>
34: e3a00000    mov    r0, #0
38: e28dd008    add    sp, sp, #8
3c: e12ffff1e    bx     lr
40: 1dcd64ff    .word 0x1dcd64ff
```

**The loop remains when we use volatile.**

# What is 'bare metal'?

The default build process for C assumes a *hosted* environment. It provides standard libraries, all the stuff that happens before `main`.

To build bare-metal, our makefile disables these defaults; we must supply our own versions when needed.

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

# Makefile settings

Compile freestanding

```
CFLAGS = -ffreestanding
```

Link without standard libs and start files

```
LDFLAGS = -nostdlib
```

Link with gcc to support division (violates

```
LDLIBS = -lgcc
```

Must supply own replacement for libs/start

That's where the fun is...!



C++ source #1 ×

A ▾



```
1 void wipe1(int arr[])
2 {
3     arr[1] = 0;
4 }
5
6 struct point {
7     int x, y, z;
8 };
9
10 void wipe2(struct point *ptr)
11 {
12     ptr->y = 0;
13 }
```

ARM gcc 5.4 (Editor #1, Compiler #1) ×

ARM gcc 5.4 ▾

-Og -ffreestanding -marm

11010

.LX0:

.text

//

Intel

A ▾



```
1 wipe1(int*):
2     mov     r3, #0
3     str     r3, [r0, #4]
4     bx      lr
5 wipe2(point*):
6     mov     r3, #0
7     str     r3, [r0, #4]
8     bx      lr
9
```

**loop:**

```
ldr r0, =0x2020001C    // set pin
str r1, [r0]
```

```
mov r2, #0x3F0000      // delay loop
```

**wait1:**

```
    subs r2, #1
    bne wait1
```

```
ldr r0, =0x20200028    // clear pin
str r1, [r0]
```

```
mov r2, #0x3F0000      // delay loop
```

**wait2:**

```
    subs r2, #1
    bne wait2
```

**b loop**

```
ldr r0, =0x2020001C
str r1, [r0]
b delay
ldr r0, =0x20200028
str r1, [r0]
b delay
b loop
```

delay:

```
mov r2, #0x3F0000
```

wait:

```
subs r2, #1
```

```
bne wait
```

// but... where to go next?

```
ldr r0, =0x2020001C
str r1, [r0]
mov r14, pc
b delay
ldr r0, =0x20200028
str r1, [r0]
mov r14, pc
b delay
b loop
```

```
delay:
mov r2, #0x3F0000
wait:
    subs r2, #1
    bne wait
mov pc, r14
```

**To return from a function, we need to know what pc we called it from, so we can continue after.**

```
ldr r0, =0x2020001C
str r1, [r0]
mov r0, #0x3F0000
mov r14, pc
b delay
ldr r0, =0x20200028
str r1, [r0]
mov r0, #0x3F0000 >> 2
mov r14, pc
b delay
b loop
```

delay:

```
subs r0, #1
```

wait:

```
bne wait
```

```
mov pc, r14
```

**We also need to pass parameters into functions.**

# Anatomy of C function call

```
int sum(int n)
{
    int total = 0;
    for (int i = 1; i < n; i++)
        total += i;
    return total;
}
```

**Call and return**

**Pass arguments**

**Local variables**

**Return value**

**Scratch/work space**

***Complication:* nested function calls, recursion**

# **Application binary interface**

**ABI specifies how code interoperates:**

- **Mechanism for call/return**
- **How parameters passed**
- **How return value communicated**
- **Use of registers (ownership/preservation)**
- **Stack management (up/down, alignment)**

**arm-none-eabi is ARM embedded ABI**  
**(“none” refers to no hosting OS)**

# Mechanics of call/return

Caller puts up to 4 arguments in r0-r3

Call instruction is **bl** (branch and link)

```
mov r0, #100
```

```
mov r1, #7
```

```
bl sum      // will set lr=pc-4
```

Callee puts return value in r0

Return instruction is **bx** (branch exchange)

```
add r0, r0, r1
```

```
bx lr      // pc=lr
```

btw: lr is mnemonic for r14



# Caller and Callee

***caller*** - function doing the calling

***callee*** - function called

main is caller of one

one is callee of main

+ caller of two

```
void main(void) {  
    add10(3);  
}
```

```
void add10(int a) {  
    add(10, a);  
}
```

```
int add(int x, int y) {  
    return x + y;  
}
```

# Register Ownership

**r0-r3** are **callee-owned** registers

- **Callee** can change these registers
- **Caller** cedes to callee, cannot assume value will be preserved across call to callee

**r4-r13** are **caller-owned** registers

- **Callee** must preserve values in these registers
- **Caller** retains ownership, expects value to be same after call as it was before call

# **Discuss**

- 1. If a function needs scratch space for an intermediate value, which type of register should it choose?**
- 2. What must a function do when it wants to use a caller-owned register?**
- 3. What is the advantage in having some registers callee-owned and others caller-owned? Why not treat all same?**
- 4. How can we implement nested calls when we only have a single shared lr register?**

# The Stack to the Rescue



# Program Stack

**Region in memory to store local variables, scratch space, save register values**


- **LIFO: push adds value on top of stack, pop removes lastmost value**
- **r13 (alias sp) points to end of stack**
- **stack grows down**
  - **newer values at lower addresses**
  - **push subtracts from sp**
  - **pop adds to sp**
- **push/pop are aliases for a general instruction (load/store multiple with writeback)**

```
// start.s
mov sp, #0x8000000
bl main
```

```
// main.c
void main(void)
{
    add10(3);
}

int add10(int a)
{
    int arr[100];
    return add(arr, 100);
}
```


*Not to scale*



*Diagram not to scale*

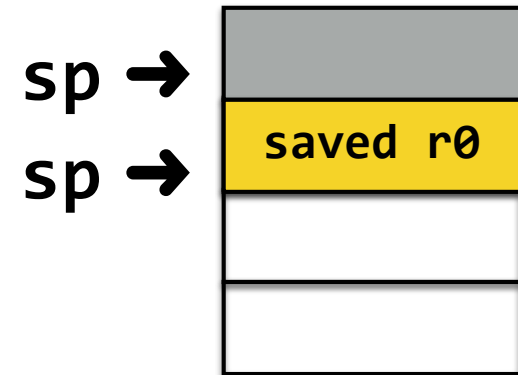
# Single stack frame

```
int add10(int a, int b)
{
    int c = 2*a;
    ...
    return c;
}
```



# Stack operations

```
// PUSH (store reg to stack)
// *-sp = r0
// decrement sp before store
push {r0}
// equivalent to:
    str r0, [sp, #-4]!
```



```
// POP (restore reg from stack)
// r0 = *sp++
// increment sp after load
pop {r0}
// equivalent to:
    ldr r0, [sp], #4
```

**“Full Descending” stack**



```
int f(int a, int b)
{
    int c = g(a);
    return b + c;
}
```

**If f calls g...**

**Why do they collide on use of `lr` ?**

**Is there similar collision for `r0`? `r1`?**

**What do we do about it?**

**use stack as temp storage!**

**example.c**

**r0**  
**r1**  
**r2**  
**r3**  
**lr**  
**sp**  
**pc**

**0x80000000**  
**0x7fffffff**  
**0x7fffffff8**  
**0x7fffffff4**  
**0x7ffffff0**  
**0x7ffffec**  
**0x7ffffe8**  
**0x7ffffe4**  
**0x7ffffe0**




	0x80000000	
	0x7fffffff	
	0x7fffffff8	
	0x7fffffff4	
	0x7ffffff0	
	0x7ffffffc	
	0x7ffffffe8	
	0x7ffffffe4	
r0	0x7ffffffe0	
r1	:	:
r2	0x7ffffe70	
r3	0x7ffffe6c	
lr	0x7ffffe68	
sp	0x7ffffe64	
	0x7ffffe60	
pc	0x7ffffe5c	
	0x7ffffe58	
	0x7ffffe54	

# sp in constant motion

Access values on stack using  
sp-relative addressing, but ....


sp is constantly changing!  
(push, pop, add sp, sub sp)



# Add frame pointer (fp)

Dedicate fp register to be used as fixed anchor

Offsets relative to fp stay constant!



# **APCS “full frame”**

**APCS = ARM Procedure Call Standard**

**Conventions for use of frame pointer + frame layout that allows for reliable stack introspection**

**gcc CFLAGS to enable: -mapcs-frame**

**r12 used as fp**

**Adds a prologue/epilogue to each function that sets up/tears down the standard frame and manages fp**

**You'll use this in assignment 4**

# Trace APCs

## *Prolog*

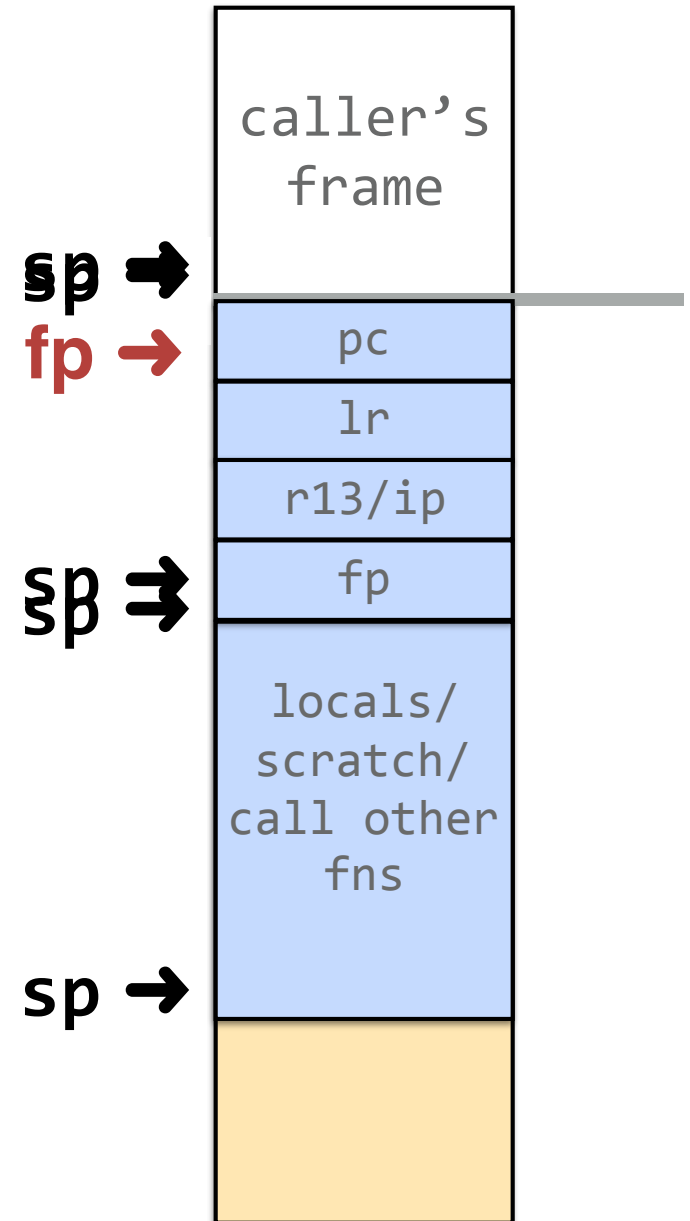
stmfd fp, sp, lr, pc  
set fp to first word of stack frame

## *Body*

fp stays anchored  
access data on stack fp-relative  
offsets won't vary even if sp changing

## *Epilog*


ldmia fp, sp, lr  
can't pop pc (**why not?**), manually adjust stack





# FPs form linked chain

*other* =  
additional saved regs,  
locals,  
scratch



```
// start.s
```

```
// Need to initialize fp = NULL  
// to terminate end of chain
```

```
    mov sp, #0x80000000  
    mov fp, #0      // fp = NULL  
    bl main
```

# APCS Pros/Cons

- + Anchored fp, offsets are constant
- + Standardized frame layout enables introspection
- + Backtrace for debugging
- + Unwind stack on exception
- Expensive, every function call affected
  - prolog/epilog add ~5 instructions
  - 4 registers push/pop => add 16 bytes per frame
  - consumes one of our precious registers